

Modeling Normative Aspects of Social Attitudes with Beta Regression

Mark H. White II¹

¹ University of Kansas

Author Note

Author note will go here.

Correspondence concerning this article should be addressed to Mark H. White II.

E-mail: markhwhiteii@gmail.com

Abstract

Abstract will go here.

Keywords: beta regression, hurdle models, norms, social attitudes

Modeling Normative Aspects of Social Attitudes with Beta Regression

Statistical Background**The Beta Distribution**

The beta distribution can be used to model the residuals in a generalized linear model when the values of the dependent variable are bounded $0 < y < 1$ (Coxe, West, & Aiken, 2013). The probability density function (pdf) of the beta distribution is determined by two parameters, α and β , that are called “shape” parameters:

$$f(y; \alpha, \beta) = \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} y^{\alpha-1} (1-y)^{\beta-1}$$

where $\Gamma(\cdot)$ is the gamma function. The two shape parameters pull the mean toward zero and one, respectively. If α is larger than β , the mean leans toward zero; if the reverse is true, the mean leans toward one. One of the benefits of the beta distribution is that it is flexible and can take a number of distributional shapes. ADD FIGURE 1 HERE TO SHOW VARIOUS SHAPES.

The values of these parameters are not meaningful to applied researchers, however. Rigby, Stasinopoulos, Heller, and De Bastiani (2017) “reparameterized” the beta distribution to make it easier understand in a regression framework (but see Ferrari & Cribari-Neto, 2004 for a different parameterization). Instead of predicting α and β , they parameterize the beta distribution with two different parameters: μ (called the “location” parameter) and σ (called the “scale” parameter), where $\mu = \frac{\alpha}{(\alpha+\beta)}$ and $\sigma = \sqrt{\frac{1}{(\alpha+\beta+1)}}$. μ is equivalent to the mean, and σ is related positively to the variance. (Note that σ is *not* the standard deviation, even though σ refers to the standard deviation in many other contexts.) The variance is equivalent to $\sigma^2 \mu(1-\mu)$. There are two important things to note from this equation: First, the greater the σ , the greater the variance; second, the variance depends on the mean, which means that beta regression, covered shortly, will be naturally heteroskedastic.

Using this distribution in a regression framework cannot handle the values zero and

one, however. How can we model dependent variables that also include observations at zero and one?

Zero-One Inflated Beta Distribution

Rigby et al. (2017) show that the beta distribution can be “inflated” at zero or one—that is, zeros and ones can be included in the distribution. This distribution can be thought of as a mixture distribution. When the dependent variable contains zeroes and ones (i.e., $0 \leq y \leq 1$), the pdf for this beta mixture, beinf , is:

$$\text{beinf}(y; \mu, \sigma, \nu, \tau) = \begin{cases} p_0 & \text{if } y = 0 \\ (1 - p_0 - p_1)f(y; \mu, \sigma) & \text{if } 0 < y < 1 \\ p_1 & \text{if } y = 1 \end{cases}$$

for $0 \leq y \leq 1$, where $f(y; \mu, \sigma)$ is the beta pdf with μ and σ bounded *between* zero and one. The two additional parameters, ν and τ , are mixture parameters. p_0 is the probability that a case equals zero, p_1 is the probability that a case equals one, and p_2 (i.e., $1 - p_0 - p_1$) is the probability that the case comes from the beta distribution. In terms of these two additional parameters, $p_0 = \frac{\nu}{(1+\nu+\tau)}$ and $p_1 = \frac{\tau}{(1+\nu+\tau)}$. Rearranging these algebraically, ν is the odds that a case is zero to being from the beta distribution, $\nu = p_0/p_2$, and τ is the odds that a case is a one to being from the beta distribution, $\tau = p_1/p_2$.

Beta Regression Models

The goal now is to predict these four parameters, μ, σ, ν, τ , from any number of predictor variables. All four parameters can be predicted by an identical set of predictors, none of the same predictors, or anywhere in between. Both μ and σ have to be between zero and one, so we can use the logistic link function to fit predicted values in this range; both ν and τ have to be greater than zero, so we can use the log link function to fit predicted values

in this range. Imagine we have one predictor, x . We could include this variable as a predictor of all four variables with the equations:

$$\log\left(\frac{\mu}{1-\mu}\right) = \beta_{10} + \beta_{11}X$$

$$\log\left(\frac{\sigma}{1-\sigma}\right) = \beta_{20} + \beta_{21}X$$

$$\log(\nu) = \beta_{30} + \beta_{31}X$$

$$\log(\tau) = \beta_{40} + \beta_{41}X$$

Or, equivalently stated:

$$\mu = \frac{1}{1+e^{-(\beta_{10}+\beta_{11}X)}}$$

$$\sigma = \frac{1}{1+e^{-(\beta_{20}+\beta_{21}X)}}$$

$$\nu = e^{\beta_{30}+\beta_{31}X}$$

$$\tau = e^{\beta_{40}+\beta_{41}X}$$

where e^x is the natural exponential function. A regression model can also be used when the dependent variable contains zeros but no ones (e.g., $0 \leq y < 1$) or when it contains ones but no zeros (e.g., $0 < y \leq 1$). Let c be the value—0 or 1—that is included. The pdf is:

$$\text{beinf}_c(y; \mu, \sigma, \nu) = \begin{cases} p_c & \text{if } y = c \\ (1 - p_c)f(y; \mu, \sigma) & \text{if } 0 < y < 1 \end{cases}$$

where $\nu = p_c/(1 - p_c)$ and everything else is the same as above. The same link functions are used as above, but the fourth parameter, τ , is not included, as the dependent variable can only take on values of c or from the beta distribution. Lastly, if no zeros or ones are observed, the beta distribution alone can be used as the pdf. This results in a beta regression where the researcher is only predicting the location, μ , and shape, σ . Since there is no mixture with values of being 0 or 1, the latter two mixture parameters, ν and τ , are not included.

Modeling Bounded Variables Outside the Unit Interval

These beta regression models are often used when dealing with rates and proportions, given that these are naturally bounded $0 \leq y \leq 1$ (Buntaine, 2011; Eskelson, Madsen, Hagar, & Temesgen, 2011; Gallardo, Bovea, Colomer, & Prades, 2012; Hubben et al., 2008; Peplonska et al., 2012). The beta distribution is doubly bounded continuous distribution, meaning that, although it is on a continuous scale, values cannot be greater than the upper bound, u , or lesser than the lower bound, l . In the case of the beta distribution, $l = 0$ and $u = 1$.

One can view many scales researchers use as doubly bounded and continuous. Although researchers generally model variables on Likert scales and sliding scales (e.g., feeling thermometers) as being conditionally normally distributed, these variables are, by definition, not strictly normal. Observations from a normal distribution can take on any value on the real number line. Observations from a standard 7-point Likert scale can only take on values 1 through 7. In studying controversial issues like prejudice and politics, many participants score at the lower or upper bounds, causing floor or ceiling effects; in these cases, the assumptions of normality and homoskedasticity are likely to be violated. One can explicitly take into account that the response is doubly bounded and heteroskedastic by using the beta regression models described above, using a simple linear rescaling of the data.

If one observes a dependent variable limited between two bounds, there is a straightforward way to rescale the variable to the $0 \leq y \leq 1$ range:

$$y'_i = (y_i - l)/(u - l)$$

where y is the variable on the original scale, y' is the rescaled variable, u is the upper bound (i.e., the largest possible value on the scale), l is the lower bound (i.e., the smallest possible value on the scale), and the i subscript denotes an individual's score. On a standard 7-point Likert scale, $l = 1$ and $u = 7$. This rescaling allows a researcher to explicitly model conditional variance, floor effects, and ceiling effects using beta regression.

Conditional variance. Having an outcome that is a doubly bounded continuous variable can produce heteroskedasticity. One of the assumptions of an ordinary least squares (OLS) regression is homoskedasticity—that the variance of the errors are unrelated to any predictor or any linear combination of predictors. A regression equation with one independent variable x is often written as $y_i = \beta_0 + \beta_1 x_i + \epsilon_i$, where ϵ_i is how far one's observed value y_i is from their predicted value (i.e., the residual). Let \hat{y}_i be the predicted value for observation i , where $\hat{y}_i = \beta_0 + \beta_1 x_i$. We can simplify the equation to: $y_i = \hat{y}_i + \epsilon_i$ where the residuals ϵ are normally distributed with a mean of zero and some variance—that is, $\epsilon \sim N(0, \sigma_\epsilon^2)$. We can further simplify this to: $y_i | x_i \sim N(\hat{y}_i, \sigma_\epsilon^2)$, which means that each observation of the dependent variable we make, given the predictors we use, are normally distributed with a mean of that individual's predicted score and some variance. The assumption of homoskedasticity is found in that σ_ϵ^2 does *not* have a subscript i . This means that there is one *common variance* of the residuals. Imagine one runs a regression and observes $\hat{y}_i = 2.5 + 1.5 \times x_i$ and $\sigma_\epsilon^2 = 9$. For the first individual in the data, say we observe $x_1 = 0$, while we observe $x_2 = 5$ for the second individual. The model then assumes that the first person comes from a normal distribution with a mean of 4 (i.e., \hat{y}_1) and a variance of 9 (i.e., σ_ϵ^2), while the second person comes from a normal distribution with a mean of 10 (i.e., \hat{y}_2) and that same variance of 9 (i.e., σ_ϵ^2). The assumption is violated when these variances are not constant across all levels of the predictor variables (or any linear combination of the predictor variables). The violation of this assumption can cause Type I or Type II errors, depending on the type of heteroskedasticity and the source [CITE].

There are a number of ways to test for and address violations of this assumption [CITE]. Many common ways of dealing with heteroskedasticity treat it as a nuisance to be dealt with, such as robust standard errors, a transformation of the dependent variable, or weighted least squares [CITE]. Heteroskedasticity can come from various sources [CITE], one being that the variance is influenced by one of the observed predictor variables. Instead of treating this as a nuisance to be corrected for when calculating p -values, explicitly modeling

this conditional variance could be an interesting phenomenon per se. Since the beta regression models predict a shape parameter, σ , a researcher can draw conclusions like: “As x increases, the variance in y increases.” Since we generally focus on the *location* parameter in OLS regression, researchers are essentially ignoring interesting, potentially theoretically-relevant parts of their data.

Floor and ceiling effects. Floor and ceiling effects occur when there are many observations at the lower and upper bound of the scale, respectively [CITE]. CONSEQUENCES OF USING OLS WITH THESE EFFECTS [CITE]. Variables displaying these effects can be thought of as being censored or truncated; a common regression technique for censored or truncated variables are Tobit models (Smithson & Merkle, 2013). Censoring occurs when all scores beyond some threshold are recorded at that threshold. If a researcher is measuring reaction times, censoring occurs if, for example, the researcher decides that any times above five seconds are scored as five. If many people score above five, this censoring can cause a ceiling effect. Truncating occurs when people who score beyond some threshold are excluded from the data. If the researcher simply does not record trials that take longer than five seconds, then the data are considered to be truncated. Tobit models can be useful in these situations. However, Tobit models assume that the underlying dimension is still normal. A good application of this is in measuring the abilities of intelligent children (McBee, 2010). The test may be too easy for the population, resulting in a ceiling effect. Nonetheless, the ability being measured is assumed to be normally distributed.

In the realm of prejudice and politics, I argue that it is useful to *not* think of prejudice as coming from a latent normal distribution. Instead, people engage in a decision-making process. Using prejudice as an example, people decide whether or not they are going to admit to feeling any prejudice (i.e., scoring at the floor or above the floor); if they decide to admit any prejudice, then they decide how much to report. As another example: When measuring attitudes toward a polarizing political figure, people decide whether or not to respond entirely in the negative (i.e., at the floor), entirely in the positive (i.e., at the

ceiling), or somewhere in between. This could result in a bimodal distribution, with both ceiling and floor effects.

This decision making process can be modeled with the inflated beta regression models described above (at zero and one, zero only, or one only). While these are often referred to as “zero-one inflated beta regression” models, a more general name used for these types of models are “two part” models. When one assumes a decision-making process behind the two parts, they are often referred to as “hurdle” models in economics [CITE].

A classic example of a hurdle model is cigarette consumption [CITE]. A researcher might be interested in how income relates to number of cigarettes purchased in a week. The researcher will observe a floor effect at zero, since many people do not smoke at all. This is a two-part process: People either decide to smoke or not, and then they decide how many cigarettes to buy. Using a hurdle model, the researcher can predict the odds of someone purchasing any cigarettes from income and then how many cigarettes someone purchases from income. We might expect that income does not influence whether or not one smokes, but could influence how many cigarettes one decides to buy.

There are a number of flavors of hurdle models. The inflated beta regression models above can be seen as sequential, independent hurdle models. That is, the estimated coefficients predicting the mixture parameters, ν and τ , are assumed to be *independent* of the location and scale parameters, μ and σ . Dependencies between them can be modeled [CITE], but I will not demonstrate these models here.

Norms Produce Invariance and Ceiling and Floor Effects

Data from MTurk

Implementation in GAMLSS

Show example from FoRS data

References

- Buntaine, M. T. (2011). Does the Asian Development Bank respond to past environmental performance when allocating environmentally risky financing? *World Development*, 39(3), 336–350.
- Coxe, S., West, S. G., & Aiken, L. S. (2013). Generalized linear models. In T. D. Little (Ed.), *The Oxford handbook of quantitative methods, volume 2* (pp. 26–51). New York, NY: Oxford University Press.
- Eskelson, B. N., Madsen, L., Hagar, J. C., & Temesgen, H. (2011). Estimating riparian understory vegetation cover with beta regression and copula models. *Forest Science*, 57(3), 212–221.
- Ferrari, S., & Cribari-Neto, F. (2004). Beta regression for modelling rates and proportions. *Journal of Applied Statistics*, 31(7), 799–815.
- Gallardo, A., Bovea, M. D., Colomer, F. J., & Prades, M. (2012). Analysis of collection systems for sorted household waste in Spain. *Waste Management*, 32(9), 1623–1633.
- Hubben, G. A. A., Bishai, D., Pechlivanoglou, P., Cattelan, A. M., Grisetti, R., Facchin, C., ... Tramarin, A. (2008). The societal burden of HIV/AIDS in Northern Italy: An analysis of costs and quality of life. *AIDS Care*, 20(4), 449–455.
- McBee, M. (2010). Modeling outcomes with floor or ceiling effects: An introduction to the tobit model. *Gifted Child Quarterly*, 54(4), 314–320.
- Peplonska, B., Bukowska, A., Sobala, W., Reszka, E., Gromadzińska, J., Wasowicz, W., ... Ursin, G. (2012). Rotating night shift work and mammographic density. *Cancer, Epidemiology, Biomarkers, & Prevention*, 21(7), 1028–1037.
- Rigby, R., Stasinopoulos, M., Heller, G., & De Bastiani, F. (2017). *Distributions for modelling location, scale, and shape: Using GAMLSS in R*. Unpublished draft. Retrieved from gamlss.com/books-articles
- Smithson, M., & Merkle, E. C. (2013). *Generalized linear models for categorical and continuous limited dependent variables*. Boca Raton, FL: CRC Press.