

Using Beta Regression to Model Normative Aspects of Prejudice and Political Attitudes

Mark H. White II¹

¹ University of Kansas

Author Note

I would like to thank Dr. Achim Zeileis and Dr. Jonathan Templin for comments on an earlier version of this manuscript. All code and data necessary to reproduce this manuscript can be accessed at https://github.com/markhwhiteii/beta_hurdle.

Correspondence concerning this article should be addressed to Mark H. White II.
E-mail: markhwhiteii@gmail.com

Abstract

Many important phenomena in the psychological study of prejudice and political attitudes involve measuring dependent variables that are strongly influenced by normative pressures. Norms produce invariance, which can create outcome variables with ceiling effects, floor effects, and heteroskedasticity. These characteristics cause problems for ordinary least squares (OLS) regression and analysis of variance (ANOVA), which assume conditional normality and constant variance. I describe the beta distribution and propose beta regression as a way to model doubly-bounded continuous variables (e.g., Likert scales, feeling thermometers). I also demonstrate how beta regression can model conditional variance, ceiling effects, and floor effects as a proxy for norms. Lastly, I provide a brief tutorial on how to implement various beta regression models using R.

Keywords: beta regression, hurdle models, gamlss, norms, social attitudes

Using Beta Regression to Model Normative Aspects of Prejudice and Political Attitudes

Behavior is a function of the person and the environment (Lewin, 1951). It is important that researchers study the effects of personality *and* situational variables, as well as how they interact. However, Louis, Mavor, and Terry (2003) argue that “normative influence in social attitudes and behavior is systematically underestimated” (p. 190). The logic behind this argument is that researchers use statistical procedures that analyze *variance* and assume normality. But norms produce *invariance*—they lead people to engage in the same behaviors or respond to survey items in the same way. Norms can produce skewed distributions, as well: floor effects (strong norm in disagreement with a survey item), ceiling effects (strong norm for agreement), or strongly leptokurtic distributions (strong norm for answering the midpoint).

Louis and colleagues note that using only the limited number of common statistical procedures (e.g., ordinary least squares, OLS, regression) influences theoretical and methodological choices: Researchers look for measures that are normally distributed, which they argue favors personality explanations and “privileges measures that tap unclear or weak norms” (p. 190). This systematically ignores phenomena that are subject to strong normative influences, which overestimates the influence of personality on attitudes and behavior.

When personality researchers *do* model these types of variables using OLS regression, it is the deviant cases that drive the linear relationships. For example, Louis et al. (2003) examine the bivariate relationship between right-wing authoritarianism (RWA) and how much people agree with the item, “The White race is the best race.” 72% of the sample selected “1, *Strongly Disagree*” to the White race item. However, what often gets reported in articles is just the correlation: $r = .34, p < .001$. We learn that RWA positively predicts White supremacist attitudes. But what are we missing when only this is reported?

Louis et al. note that, for 72% of the sample, a correlation can not even be calculated—there was no variance in the scores. Everyone in this 72% adhered to the egalitarian norm that it is not OK to openly say the White race is the best race. For the rest

of the sample, the correlation was .19. The authors raise two concerns: first, that researchers could decide to exclude items like this because they aren't normally distributed, which prevents one from studying an important construct (White supremacy); second, readers interpret correlations symmetrically. A "symmetric" interpretation is that people at *both* the high and low ends of the authoritarianism spectrum are driving the relationship, when we know that is only the top 28% of the sample.

Aims of the Present Paper

I second the concerns raised by Louis et al. (2003). The purpose of this paper is to propose a concrete solution to the issues raised. They lamented that there is "no conventional analysis of invariance" (p. 192), and they gave no suggestions for how else to analyze data to mitigate these issues. I propose beta regression (presented a year after Louis and colleagues' paper) as one solution. Smithson and Verkuilen (2006) touch on this briefly, but I fully extend the argument here.

I have three goals in the present paper: first, to familiarize researchers with the beta distribution and beta regression, which can handle non-normal data and model conditional variance; second, to empirically demonstrate Louis et al.'s contention that norms produce invariance; third, to give a brief tutorial on how researchers can use beta regression to address the crucial issues raised by Louis et al. (2003), such as modeling normative aspects in survey responses.

Statistical Background

The Beta Distribution

The beta distribution can be used in a generalized linear model when the values of the dependent variable are bounded $0 < y < 1$ (Coxe, West, & Aiken, 2013). The probability density function (pdf) of the beta distribution is determined by two parameters, α and β , that are called "shape" parameters:

$$f(y; \alpha, \beta) = \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} y^{\alpha-1} (1-y)^{\beta-1}$$

where $\Gamma(\cdot)$ is the gamma function. One of the benefits of the beta distribution is that it is flexible and can take a number of distributional shapes (Figure 1).

These parameters are not inherently meaningful to researchers. Rigby, Stasinopoulos, Heller, and De Bastiani (2017) “reparameterized” the beta distribution so that the two parameters determining the shape of the distribution would be more useful in a regression framework (see Ferrari & Cribari-Neto, 2004 for a different parameterization). Instead of predicting α and β , they parameterize (i.e., algebraically rearrange parameters) the distribution so that beta regression predicts two different parameters: μ (called the “location” parameter) and σ (called the “scale” parameter), where $\mu = \frac{\alpha}{\alpha+\beta}$ and $\sigma = \sqrt{\frac{1}{(\alpha+\beta+1)}}$. μ is equivalent to the mean, and σ is related positively to the variance (note that σ is *not* the standard deviation, even though σ commonly refers to the standard deviation). The variance is equivalent to $\sigma^2 \mu(1-\mu)$. There are two things to note from this equation: first, the greater the σ , the greater the variance; second, the variance depends on the mean, which means that beta regression will be naturally heteroskedastic (covered shortly). Using this distribution in a regression framework cannot handle the dependent variable taking on values 0 and 1, however. How can we model dependent variables in the range $0 \leq y \leq 1$?

Zero-One Inflated Beta Distribution

Rigby et al. (2017) show that the beta distribution can be “inflated” at 0 and 1, meaning the dependent variable contains 0s and 1s (i.e., $0 \leq y \leq 1$), the pdf for this beta inflated distribution, beinf , is:

$$\text{beinf}(y; \mu, \sigma, \nu, \tau) = \begin{cases} p_0 & \text{if } y = 0 \\ (1 - p_0 - p_1)f(y; \mu, \sigma) & \text{if } 0 < y < 1 \\ p_1 & \text{if } y = 1 \end{cases}$$

for $0 \leq y \leq 1$, where $f(y; \mu, \sigma)$ is the beta pdf with μ and σ bounded *between* 0 and 1. The two additional parameters, ν and τ , are related to p_0 and p_1 , respectively. p_0 is the probability that a case equals 0, p_1 is the probability that a case equals 1, and p_2 (i.e., $1 - p_0 - p_1$) is the probability that the case comes from the beta distribution. In terms of these two additional parameters, $p_0 = \frac{\nu}{(1+\nu+\tau)}$ and $p_1 = \frac{\tau}{(1+\nu+\tau)}$. Rearranging these algebraically, ν is the odds that a case is 0 to being $0 < y < 1$, $\nu = p_0/p_2$, and τ is the odds that a case is a 1 to being $0 < y < 1$, $\tau = p_1/p_2$. It should be noted, however, that this distribution only includes *one* source of 0 values (i.e., p_0). For this reason, some would not call it an “inflated” distribution—which has more than one source of 0 values—but instead a “hurdle model” (discussed below, see Cox et al., 2013).

Beta Regression Models

The goal for a researcher is to predict these four parameters, μ, σ, ν, τ , from variables of theoretical interest. All four parameters can be predicted by an identical set of predictors, none of the same predictors, or any combination. Both μ and σ have to be between 0 and 1 (since the beta distribution is bounded between 0 and 1), so one can use the logistic link function to fit predicted values in this range. Both ν and τ have to be greater than 0 (since they are odds), so one can use the log link function to fit predicted values in this range. Imagine a researcher has one predictor of interest, x . They could include this variable as a predictor of all four parameters with the following equations:

$$\log\left(\frac{\mu}{1-\mu}\right) = \beta_{10} + \beta_{11}X$$

$$\log\left(\frac{\sigma}{1-\sigma}\right) = \beta_{20} + \beta_{21}X$$

$$\log(\nu) = \beta_{30} + \beta_{31}X$$

$$\log(\tau) = \beta_{40} + \beta_{41}X$$

Or equivalently:

$$\mu = \frac{1}{1+e^{-(\beta_{10}+\beta_{11}X)}}$$

$$\begin{aligned}\sigma &= \frac{1}{1+e^{-(\beta_{20}+\beta_{21}X)}} \\ \nu &= e^{\beta_{30}+\beta_{31}X} \\ \tau &= e^{\beta_{40}+\beta_{41}X}\end{aligned}$$

where e^x is the natural exponential function. A different inflated beta distribution can be used when the dependent variable contains 0s but not 1s (e.g., $0 \leq y < 1$) or when it contains 1s but not 0s (e.g., $0 < y \leq 1$). Let c be the value—0 or 1—that is included (i.e., inflated). The pdf is:

$$\text{beinf}_c(y; \mu, \sigma, \nu) = \begin{cases} p_c & \text{if } y = c \\ (1 - p_c)f(y; \mu, \sigma) & \text{if } 0 < y < 1 \end{cases}$$

where $\nu = p_c/(1 - p_c)$ and the remaining notation is the same as in the zero-and-one inflated beta distribution. The same link functions are used, but the fourth parameter, τ , is not included, since the distribution is only inflated at c . Lastly, if no 0s or 1s are observed, the beta distribution alone can be used as the pdf. This results in a beta regression where the researcher is only predicting the location, μ , and shape, σ . Since there is no inflation at 0 or 1, the latter two parameters, ν and τ , are not included.

Modeling Bounded Variables Beyond the Zero Through One Range

Beta regression is commonly used when rates and proportions are dependent variables, given that these are naturally bounded $0 \leq y \leq 1$ (Buntaine, 2011; Eskelson, Madsen, Hagar, & Temesgen, 2011; Gallardo, Bovea, Colomer, & Prades, 2012; Hubben et al., 2008; Peplonska et al., 2012). More generally, the beta distribution is a doubly-bounded continuous distribution: Although it is on a continuous scale, values cannot exceed the upper bound, u , or be less than the lower bound, l . In the case of the zero-and-one inflated beta distribution, $l = 0$ and $u = 1$.

Most scales used to measure prejudice and political attitudes are doubly bounded and continuous. Although researchers generally model variables on Likert scales and sliding scales (e.g., feeling thermometers) as being conditionally normally distributed (i.e., using OLS regression), these variables are not strictly normal. Observations from a normal distribution can take on any value on the real number line, while observations from a standard 7-point Likert scale can only take on values 1 through 7. In studying controversial and polarizing issues like prejudice and politics, many participants score at the lower or upper bounds, causing floor or ceiling effects. In these cases, the assumptions of normality and homoskedasticity are likely to be violated. Many times researchers create, update, or seek out measures to so that the distributional form of the dependent variable takes on a normal shape—statistical assumptions are directing the phenomenon researchers choose to observe and the theoretical questions that they ask. Instead, one can explicitly take into account that the response is doubly-bounded and heteroskedastic by using the beta regression models described above, using a simple linear rescaling of the data.

If one observes a dependent variable limited between two bounds, there is a straightforward way to rescale the variable to the $0 \leq y \leq 1$ range:

$$y'_i = (y_i - l)/(u - l)$$

where y is the variable on the original scale, y' is the rescaled variable, u is the upper bound (i.e., the largest possible value on the scale), l is the lower bound (i.e., the smallest possible value on the scale), and the i subscript denotes an individual's score. On a standard 7-point Likert scale, $l = 1$ and $u = 7$. This rescaling allows a researcher to explicitly model conditional variance, floor effects, and ceiling effects using beta regression.

Conditional variance. Observing a dependent variable that is doubly-bounded and continuous can produce heteroskedasticity, which violates the assumption of homoskedasticity in OLS regression. This assumption holds that the variance of the errors is unrelated to any predictor, any linear combination of predictors, and the predicted values. A regression equation with one predictor variable x is often written as $y_i = \beta_0 + \beta_1 x_i + \epsilon_i$,

where ϵ_i is how far an observed value y_i is from the predicted value (i.e., the residual). Let \hat{y}_i be the predicted value for observation i , where $\hat{y}_i = \beta_0 + \beta_1 x_i$. We can simplify the equation to $y_i = \hat{y}_i + \epsilon_i$, where the residuals ϵ are normally distributed with a mean of 0 and some variance—that is, $\epsilon \sim N(0, \sigma_\epsilon^2)$. We can further simplify this to: $y_i | x_i \sim N(\hat{y}_i, \sigma_\epsilon^2)$, which means that each observation of the dependent variable we make, given the predictors we use, is normally distributed with a mean of that observation’s predicted score and some variance. The assumption of homoskedasticity is found in that σ_ϵ^2 does *not* have a subscript i . This means that there is one *common variance* of the residuals.

Imagine one runs a regression and observes $\hat{y}_i = 2.5 + 1.5 \times x_i$ and $\sigma_\epsilon^2 = 9$. Say that the first individual’s score on x is 0, $x_1 = 0$, and the second individual’s score is $x_2 = 5$. The model then assumes that the first person comes from a normal distribution with a mean of 4 (i.e., \hat{y}_1) and a variance of 9 (i.e., σ_ϵ^2), while the second person comes from a normal distribution with a mean of 10 and that same variance of 9. The violation of this assumption can lead to inflated Type I errors or diminished statistical power, depending on the type and source of the heteroskedasticity (Hayes & Cai, 2007; Long & Ervin, 2000; Rosopa, Schaffer, & Schroeder, 2013).

While common remedies for heteroskedasticity (e.g., robust standard errors, transforming the dependent variable, weighted least squares) treat it as a nuisance to be corrected for when calculating p -values (Hayes & Cai, 2007; Long & Ervin, 2000; Rosopa et al., 2013), it could be that explicitly modeling this conditional variance is an interesting phenomenon itself, with important theoretical implications. Heteroskedasticity can arise when the error variance is predicted by one of the observed predictor variables, and since beta regression models predict a shape parameter, σ , a researcher can explicitly model this relationship. This allows researchers to ask questions about conditional variance: “Does the variance in y increase as x increases?” Since the focus of OLS regression is the *location* parameter (e.g., conditional mean), researchers might be ignoring interesting, potentially theoretically-relevant aspects of their data. After discussing floor and ceiling effects, I will

argue that normative influences are one of these aspects.

Floor and ceiling effects. Floor and ceiling effects occur when there are many observations of the dependent variable at the lower and upper bound of the scale, respectively (Everitt & Skrondal, 2010). Variables displaying these effects are likely to violate numerous assumptions of OLS regression, such as conditional normality and homoskedasticity. One option is to model them as censored or truncated. Censoring occurs when all scores beyond some threshold are recorded at that threshold. For example, if a researcher is measuring reaction times, censoring occurs if they decide that any times above 5 seconds are scored as 5. If many people took longer than 5 seconds, this censoring can cause a ceiling effect. Truncating occurs when people who score beyond some threshold are excluded from the data. If the researcher simply does not record trials that take longer than 5 seconds, then the data are truncated. A common regression technique for censored or truncated variables are Tobit models (McBee, 2010; Smithson & Merkle, 2013). However, standard Tobit models assume that the underlying latent construct is normally distributed.

In the realm of prejudice and politics, I do *not* believe that it is useful to think of prejudice as coming from a latent normal distribution. Imagine a researcher observes a floor effect when measuring prejudice on a Likert scale bounded at 1 and 7. I find it unlikely that simply moving the lower bound of this scale to -7 would yield a normal distribution in responses. Instead, participants would likely opt for the new floor. This can be thought of as a decision-making process. People decide whether or not they are going to admit to feeling any prejudice. If they decide not to, they simply circle or click all 1s (i.e., the floor). If they decide to admit prejudice, they respond to the items. As another example: When offering one's attitudes toward a polarizing political figure, a participant decides whether or not to respond entirely in the negative (i.e., at the floor), entirely in the positive (i.e., at the ceiling), or somewhere in between. This could result in a distribution with both ceiling and floor effects.

This decision-making process can be modeled with the inflated beta regression models

described above. The beta distribution with inflations at zero, one, or zero-and-one allow researchers to measuring ceiling effects, floor effects, or both simultaneously. While methodologists in beta regression refer to these models as “inflated” models (e.g., Abdel-Karim, 2017; Ospina & Ferrari, 2010), a more general label is “two-part” models (Coxe et al., 2013). When one assumes a decision-making process behind the two parts, these models are often referred to as “hurdle” models in the econometrics literature (Cameron & Trivedi, 2005; Wooldridge, 2010).

Cragg (1971) demonstrated that hurdle models are a generalization of the Type I Tobit model. In economics, hurdle models are often used when dependent variables are counts with excess 0s (see Carlevaro, Croissant, & Hoareau, 2017 for a review of applications). The likelihood function for these models can be separated in two parts: first, a binary regression (e.g., logistic or probit) modeling the probability of observing $y = 0$ versus $y > 0$; second, a truncated negative binomial model using just the cases where $y > 0$.

Inflated beta regressions can be seen as hurdle models, as the likelihood for the zero-and-one inflated beta regression can also be separated in two parts: first, a multinomial logistic regression modeling the probabilities of observing $y = 0$, $y = 1$, or $0 < y < 1$; second, a beta regression model using just the cases where $0 < y < 1$. Similarly, the inflated beta regression at *either zero or one* has a likelihood function that can be separated in two parts: first, a logistic regression modeling the probability of observing $y = c$ or $0 < y < 1$; second, a beta regression modeling just the cases where $0 < y < 1$ (Rigby et al., 2017).

Since their likelihoods can be separated, this implies an assumption: The two processes are independent, given the observed predictors. In other words, the error terms of each submodel are assumed to be independent of one another (where submodels refer to models predicting each parameter, μ , σ , ν , or τ , separately). Dependencies between them can be modeled (referred to as “selection” models; Carlevaro et al., 2017; Wooldridge, 2010), but I will not discuss these models here, as they are not extended to beta regression.

Norms Produce Invariance, Ceiling and Floor Effects

Louis et al. (2003) argue that norms produce invariance. I empirically demonstrate the relationship between normative strength and invariance of attitudes, replicating the methodology used by Crandall, Eshleman, and O'Brien (2002). In their Study 1, the authors presented participants with 105 social groups (e.g., Black people, librarians, deaf people, ugly people, Nazis, drug dealers). Participants were assigned to one of two groups: Half of the participants indicated how negatively they felt toward the group (i.e., self-reported prejudice), while the other half indicated how socially acceptable they perceived it was to feel and express negative feelings toward that group (i.e., acceptability of prejudice). Crandall and colleagues averaged the prejudice and acceptability scores at the group level, resulting in an $N = 105$ and two variables for each group: prejudice and acceptability. They found that expressed prejudice and acceptability of prejudice correlated at $r = .96$; people report prejudices that are socially acceptable (see Cary & Page-Gould, 2014 for a replication).

I focused on predicting the *variance* of self-reported prejudice, predicting that a stronger norms would be associated with less variance. I recruited 405 people from Amazon's Mechanical Turk website and randomly assigned them to either a *prejudice* or *acceptability* condition. All participants answered questions about 30 groups, randomly sampled from a pool of 120 groups. In the prejudice condition, participants were asked, "How do you feel toward each of these groups?" and to indicate their feelings on a scale from "very negatively (0) to very positively (100)." I reverse-scored these items so that higher scores indicate more prejudice. In the acceptability condition, participants were asked, "How OK is it in society for people to feel and say negative things about each of these groups?" and to indicate their perception on a scale from "totally unacceptable (0) to completely acceptable (100) for people to feel and say negative things about these groups."

To demonstrate the relationship between norms and invariance, I calculated the median social acceptability score and the variance of the self-reported prejudice scores for each of the 120 groups. There was no linear relationship between social acceptability and

variance of reported prejudice, $b = 0.31$, $SE = 0.84$, $t(118) = 0.37$, $p = .714$; however, there was a quadratic relationship when regressing variance of reported prejudice on social acceptability squared, $\Delta R^2 = .28$, $b = -1250.07$, $SE = 187.44$, $t(117) = -6.67$, $p < .001$ (Figure 2). The more participants perceived prejudice to be influenced by norms—either being totally unacceptable (0) or completely acceptable (100)—the lower was the variance of self-reported prejudice by other participants.

I also investigated normative influences on floor effects. Crandall (1994) proposed a measure for anti-fat prejudice; in Study 5 of this paper, Crandall makes the argument that anti-fat attitudes are more socially acceptable (and less subject to social desirability biases) than anti-Black prejudice. He does this by examining the decision-making process described above with regard to hurdle models: How often do people *always* select the least prejudiced option? That is, how often do people *always* respond at the lower bound of the distribution? He calls the proportion of people at the floor of a scale the “politically correct (PC) index,” as this shows how often people respond in the most socially acceptable (i.e., politically correct) way possible. He finds that his anti-fat scale is less subject to these floor effects: 10% of participants always responded at the lower bound for anti-Black prejudice, while this was only 3% for the anti-fat scale.

To show that norms are associated with floor effects, I calculated Crandall’s PC index for each of the 120 groups by counting how many people responded with a 0 on the reverse-scored prejudice scale (i.e., “totally positive” feelings). Using a negative binomial regression (Venables & Ripley, 2002) for these count data, I regressed the PC index on social acceptability of the prejudice. Perceived social acceptability of the prejudice by one group of participants negatively predicted the PC index, which was calculated from a *separate group* of participants, $b = -0.03$, $SE = 0.002$, $Z = -10.97$, $p < .001$ (Figure 3). The more socially acceptable a prejudice, the less people opt for the “politically correct” floor of the scale.

This evidence shows that norms do positively predict invariance. Predicting the scale parameter, σ , in beta regression could be used to investigate how much people adhere to

norms regarding the expression of certain attitudes. These data also show that normative influence can lead to floor effects (or ceiling effects). Predicting the shape parameters, ν and τ , in inflated beta regression could also be used to examine how much people adhere to norms. I now turn to demonstrating how these beta regression models can be run in R (R Core Team, 2017) via the `gamlss` package (Rigby & Stasinopoulos, 2005).

Implementation in GAMLSS

The `gamlss` name stands for generalized additive models for location (i.e., μ), scale (i.e., σ), and shape (i.e., ν and τ). This package gives researchers the ability to use a wide variety of models, including beta regression models. A number of algorithms to estimate the coefficients can be used, described in detail by Stasinopoulos, Rigby, Heller, Voudouris, and De Bastiani (2017). To first demonstrate how the package works according to the parameterization discussed in the Statistical Background section, I will generate a zero-and-one inflated beta distribution and estimate the parameters with an intercepts-only regression using the `gamlss` function. Then, I will demonstrate how to model social attitudes with real data.

Intercepts-Only Regression

We can set a number of parameters for this zero-and-one inflated distribution using the following code:

```
n <- 5000
mu <- 0.40
sigma <- 0.60
p0 <- 0.13
p1 <- 0.17
p2 <- 1 - p0 - p1
```

```
a <- mu * (1 - sigma ^ 2) / (sigma ^ 2)
b <- a * (1 - mu) / mu
```

The sample size $n = 5000$, the mean $\mu = 0.40$, the scale parameter $\sigma = 0.60$, the probability of being 0 is $p_0 = 0.13$, the probability of being 1 is $p_1 = 0.17$, and the probability of being between 0 and 1 is $p_2 = 1 - p_0 - p_1 = .70$. Using the equations in the Statistical Background section, we can convert μ and σ back to the original shape parameters a (α) and b (β). The dependent variable y can now be generated using the following code:

```
set.seed(1839)
y <- rbeta(n, a, b)
cat <- sample(1:3, n, prob = c(p0, p2, p1), replace = TRUE)
y[cat == 1] <- 0
y[cat == 3] <- 1
```

I draw n random numbers from the beta distribution with shape parameters a and b . I then draw n numbers ranging from 1 to 3 with probabilities p_0 , p_2 , and p_1 , respectively. If 1 was drawn, that observation is 0; if 2 was drawn, it remains being from the beta distribution; if 3 was drawn, that observation is 1. In this way, p_0 , p_2 , and p_1 determine the probability of being 0, from a beta distribution, or being 1, respectively.

The `gamlss` function has four arguments for formulas—one for each of the parameters. It also has an argument, `family`, that determines what distribution will be used in fitting the model. For a zero-and-one inflated beta regression, the family is `BEINF()`. One can fit the intercepts-only model using the following code:

```
fit <- gamlss(
  formula = y ~ 1,      # formula for mu
  formula.sigma = ~ 1, # formula for sigma
```

```

formula.nu = ~ 1,      # formula for nu
formula.tau = ~ 1,     # formula for tau
family = BEINF()       # distribution for model
)

```

One can use the inverse link functions to transform the coefficients back into the original scale as to compare them to the parameters set above. The estimated parameters are calculated using the following code, according to the equations and link functions described in the Statistical Background section:

```

inv_logit <- function(x) exp(x) / (1 + exp(x)) # inverse of link function
fit_mu <- inv_logit(fit$mu.coefficients)
fit_sigma <- inv_logit(fit$sigma.coefficients)
fit_nu <- exp(fit$nu.coefficients)
fit_tau <- exp(fit$tau.coefficients)
fit_p0 <- fit_nu / (1 + fit_nu + fit_tau)
fit_p1 <- fit_tau / (1 + fit_nu + fit_tau)

```

The estimates for μ , σ , p_0 , and p_1 are .406, .597, .136, and .171, respectively. These are close estimates to the population parameters set above. In research settings, each of the parameters would not be set to one number (i.e., an intercept), but be regressed on observed variables of theoretical interest. For an example, I now turn to a model using real data.

Modeling Political Ideology

The social dominance orientation (SDO) scale measures how much one supports social hierarchy and inequality in society (Ho et al., 2015; Pratto, Sidanius, Stallworth, & Malle, 1994). In personality psychology, it has a robust, oft-replicated positive relationship with political conservatism (Duriez & Van Hiel, 2002; Hiel & Mervielde, 2002; Ho et al., 2015; Pratto et al., 2000; Pratto, Stallworth, & Conway-Lanz, 1998; Pratto, Stallworth, &

Sidanius, 1997; Wilson & Sibley, 2013). Some consider SDO a central aspect of contemporary American conservatism (Eidelman, Crandall, Goodman, & Blanchar, 2012; Jost, Glaser, Kruglanski, & Sulloway, 2003). However, openly admitting that some groups are inherently better than others is a socially unacceptable thing to express, often leading to floor effects. I regress SDO on conservatism in both an OLS and beta regression framework to compare and contrast the two methods.

The data come from Study 1 of White and Crandall (2017), where 175 participants answered an 8-item, 7-point SDO scale (Ho et al., 2015) and two 7-point items that assessed conservatism (i.e., Democrat, 1, to Republican, 7, and liberal, 1, to conservative, 7). Scales were calculated by averaging their respective items. Figure 4 shows the histogram of the SDO scale—30% of the sample responded with 1s for each of the 8 items, demonstrating a floor effect.

The two variables correlated at $r = .50, p < .001$. Figure 5 shows the predicted values against the standardized residuals from regressing SDO on conservatism using OLS, with a marginal histogram showing the distribution of residuals. A loess line is drawn, and a slight jitter has been added to the points, given that many points overlapped with one another. One can see a spread of residuals getting wider as predicted values increases, suggesting heteroskedasticity; moreover, the histogram shows that the residuals are positively skewed, violating the assumption of conditional normality. How does this compare with beta regression?

First, I need to rescale the data. SDO was measured on a 1-to-7 scale, so $l = 1$ and $u = 7$. One can use the following R code to rescale a variable to $0 \leq y \leq 1$, following the equation described in the Modeling Bounded Variables Beyond the Zero Through One Range section:

```
dat$sdo <- (dat$sdo - 1) / (7 - 1)
```

It is important to remember that l and u are *not* the minimum and maximum *observed*

values, but the minimum and maximum *possible* values; for example, if lowest SDO score observed was 2, l would still be 1.

The last thing a researcher has to do is decide which value to use for the **family** argument in the **gamlss** function. Researchers can follow Table 1 for deciding which value is appropriate. One can check to see the minimum and maximum observed by using the **range()** function in R. In this case, **range(dat\$sdo)** returns 0 1, showing that both 0s and 1s are observed. According to Table 1, I should use **BEINF()** for the **family**. I run the beta regression model, using conservatism (i.e., **rw_polid**) to predict all four parameters, with the following code:

```
fit_beinf <- gamlss(sdo ~ rw_polid, ~ rw_polid, ~ rw_polid, ~ rw_polid,
                  data = dat, family = BEINF())
```

Finally, one can use the **summary()** function to examine the results. Table 2 shows the coefficients tables for each of the four submodels from the summary. The μ submodel tells the same story as the OLS model: The more conservative one is, the more they support hierarchy in society $b = 0.28, SE = 0.05, t(167) = 5.46, p < .001$. Since a logistic link function is used, however, this relationship is not linear (see Figure 7). The σ submodel shows that, as people get more conservative, they also display more variance in their responses to the SDO scale, $b = 0.14, SE = 0.06, t(167) = 2.48, p = .014$; conservatives are not just more supportive of hierarchy, they also display a wider range of attitudes toward hierarchy. In line with the evidence presented above that norms produce invariance, one explanation for this relationship is that conservatives are less influenced by norms against expressing attitudes in favor of social inequality. The conservatism coefficient for the ν model is also significant: The more conservatism one reports, the less likely they are to respond at the floor (i.e., in a “politically correct” manner; Crandall, 1994), $b = -0.54, SE = 0.12, t(167) = -4.46, p < .001$. Again, an interpretation here is that conservatives are less influenced by the normative pressure to completely reject social inequality and group hierarchies. Lastly, the coefficient

for the τ submodel is not significant; however, this could be an issue of statistical power, given that only one participant responded at the ceiling.

How do the OLS and beta regression results compare? Figure 6 shows a scatterplot of SDO regressed onto conservatism using OLS regression. The solid line is the regression line. The estimated standard deviation of the residuals, σ_e , is added to and subtracted from this line; these are the dotted lines. In line with the assumption of homoskedasticity, this standard deviation is constant across all levels of conservatism. Importantly, the lower dotted line dips below possible observed values—the model assumes that predicted values come from a distribution that has values in it that *cannot* be observed with the Likert scale used.

Figure 7 shows the same figure for the zero-and-one inflated beta regression model. The width of the dotted line increases as conservatism increases, demonstrating the heteroskedasticity of the model. This also shows that variance in SDO scores increases with conservatism; also note that the dotted line does not dip below the floor, as values in the beta distribution cannot be less than the lower bound.

Diagnostic Checks and Other Models with Conditional Variance

Diagnostic checks—such as inspecting the residuals—for beta regression models is still an active area of research, with many different authors proposing different techniques for diagnosing issues with these models (Espinheira, Ferrari, & Cribari-Neto, 2008b, 2008a; Espinheira, Santos, & Cribari-Neto, 2017; Pereira, 2017). There does not seem yet to be an agreed-upon method for performing diagnostic checks on these models.

Beta regression can also be implemented in the `betareg` package (Cribari-Neto & Zeileis, 2010; Grün, Kosmidis, & Zeileis, 2012), but it does not allow for modeling zeros and ones as easily. The `ZOIP` package (Diaz Zapata, 2017) is also available, which can be used to model mixed effects. Both of these packages use a separate parameterization than the `gamlss` package. Beta regression is not the only way to model conditional variance—I chose to focus on the method here because it is a flexible distribution that handles a large range of

non-normal shapes. Dispersion can also be predicted using double generalized linear models in the `dglm` R package (Smyth, 1989; Smyth & Verbyla, 1999) or heteroskedastic censored and truncated regression using the `crch` R package (Messner, Mayr, & Zeileis, 2016; Messner, Zeileis, Broecker, & Mayr, 2014), among others.

Conclusion

Common statistical procedures that assume normality and homoskedasticity ignore interesting theoretical phenomena in data. I described the beta regression, demonstrated that norms can produce invariance, and provided a tutorial on how to fit numerous beta-regression models using free statistical software. Louis et al. (2003) argued that common statistical procedures (a) motivate researchers to seek out measures that match the assumptions of common statistical procedures, allowing the tools we use to influence the phenomena we study, and (b) provide a misleading picture of the relationships between variables that are strongly influenced by norms. I believe beta regression can be a useful tool for addressing these issues.

References

- Abdel-Karim, A. H. (2017). Extended zero-one inflated beta and adjusted three-part regression models for proportional data analysis. *Communications in Statistics - Simulation and Computation*, 46(8), 6155–6172.
- Buntaine, M. T. (2011). Does the Asian Development Bank respond to past environmental performance when allocating environmentally risky financing? *World Development*, 39(3), 336–350.
- Cameron, A. C., & Trivedi, P. K. (2005). *Microeconometrics: Methods and applications*. New York, NY: Cambridge University Press.
- Carlevaro, F., Croissant, Y., & Hoareau, S. (2017). *Multiple hurdle Tobit models in R: The mhurdle package*. Retrieved from <https://cran.r-project.org/web/packages/mhurdle/vignettes/mhurdle.pdf>
- Cary, L. A., & Page-Gould, E. (2014). The prevalence and impact of sexism, racism and homophobia in online gaming environments. Poster presented at the annual meeting of the Society for Personality and Social Psychology, Austin, TX.
- Coxe, S., West, S. G., & Aiken, L. S. (2013). Generalized linear models. In T. D. Little (Ed.), *The Oxford handbook of quantitative methods, Volume II* (pp. 26–51). New York, NY: Oxford University Press.
- Cragg, J. G. (1971). Some statistical models for limited dependent variables with application to the demand for durable goods. *Econometrica: Journal of the Econometric Society*, 39(5), 829–844.
- Crandall, C. S. (1994). Prejudice against fat people: Ideology and self-interest. *Journal of Personality and Social Psychology*, 66(5), 882–894.
- Crandall, C. S., Eshleman, A., & O'Brien, L. (2002). Social norms and the expression and suppression of prejudice: The struggle for internalization. *Journal of Personality and Social Psychology*, 82(3), 359–378.
- Cribari-Neto, F., & Zeileis, A. (2010). Beta regression in R. *Journal of Statistical Software*,

34(2).

Diaz Zapata, J. C. (2017). *ZOIP: ZOIP distribution, ZOIP regression, ZOIP mixed regression*. Retrieved from <https://CRAN.R-project.org/package=ZOIP>

Duriez, B., & Van Hiel, A. (2002). The march of modern fascism: A comparison of social dominance orientation and authoritarianism. *Personality and Individual Differences*, 32(7), 1199–1213.

Eidelman, S., Crandall, C. S., Goodman, J. A., & Blanchar, J. C. (2012). Low-effort thought promotes political conservatism. *Personality and Social Psychology Bulletin*, 38(6), 808–820.

Eskelson, B. N., Madsen, L., Hagar, J. C., & Temesgen, H. (2011). Estimating riparian understory vegetation cover with beta regression and copula models. *Forest Science*, 57(3), 212–221.

Espinheira, P. L., Ferrari, S. L., & Cribari-Neto, F. (2008a). Influence diagnostics in beta regression. *Computational Statistics & Data Analysis*, 52(9), 4417–4431.

Espinheira, P. L., Ferrari, S. L., & Cribari-Neto, F. (2008b). On beta regression residuals. *Journal of Applied Statistics*, 35(4), 407–419.

Espinheira, P. L., Santos, E. G., & Cribari-Neto, F. (2017). On nonlinear beta regression residuals. *Biometrical Journal*, 59(3), 445–461.

Everitt, B., & Skrondal, A. (2010). *The Cambridge dictionary of statistics*. New York, NY: Cambridge University Press.

Ferrari, S. L., & Cribari-Neto, F. (2004). Beta regression for modelling rates and proportions. *Journal of Applied Statistics*, 31(7), 799–815.

Gallardo, A., Bovea, M. D., Colomer, F. J., & Prades, M. (2012). Analysis of collection systems for sorted household waste in Spain. *Waste Management*, 32(9), 1623–1633.

Grün, B., Kosmidis, I., & Zeileis, A. (2012). Extended beta regression in R: Shaken, stirred, mixed, and partitioned. *Journal of Statistical Software*, 48(11).

Hayes, A. F., & Cai, L. (2007). Using heteroskedasticity-consistent standard error estimators

- in OLS regression: An introduction and software implementation. *Behavior Research Methods*, 39(4), 709–722.
- Hiel, A. V., & Mervielde, I. (2002). Explaining conservative beliefs and political preferences: A comparison of social dominance orientation and authoritarianism. *Journal of Applied Social Psychology*, 32(5), 965–976.
- Ho, A. K., Sidanius, J., Kteily, N., Sheehy-Skeffington, J., Pratto, F., Henkel, K. E., . . . Stewart, A. L. (2015). The nature of social dominance orientation: Theorizing and measuring preferences for intergroup inequality using the new SDO₇ scale. *Journal of Personality and Social Psychology*, 109(6), 1003–1028.
- Hubben, G. A. A., Bishai, D., Pechlivanoglou, P., Cattelan, A. M., Grisetti, R., Facchin, C., . . . Tramarin, A. (2008). The societal burden of HIV/AIDS in Northern Italy: An analysis of costs and quality of life. *AIDS Care*, 20(4), 449–455.
- Jost, J. T., Glaser, J., Kruglanski, A. W., & Sulloway, F. J. (2003). Political conservatism as motivated social cognition. *Psychological Bulletin*, 129(3), 339–375.
- Lewin, K. (1951). *Field theory in social science: Selected theoretical papers*. (D. Cartwright, Ed.). New York, NY: Harper & Row.
- Long, J. S., & Ervin, L. H. (2000). Using heteroscedasticity consistent standard errors in the linear regression model. *The American Statistician*, 54(3), 217–224.
- Louis, W. R., Mavor, K. I., & Terry, D. J. (2003). Reflections on the statistical analysis of personality and norms in war, peace, and prejudice: Are deviant minorities the problem? *Analyses of Social Issues and Public Policy*, 3(1), 189–198.
- McBee, M. (2010). Modeling outcomes with floor or ceiling effects: An introduction to the tobit model. *Gifted Child Quarterly*, 54(4), 314–320.
- Messner, J. W., Mayr, G. J., & Zeileis, A. (2016). Heteroscedastic censored and truncated regression with crch. *The R Journal*, 8(1), 173–181.
- Messner, J. W., Zeileis, A., Broecker, J., & Mayr, G. J. (2014). Probabilistic wind power forecasts with an inverse power curve transformation and censored regression. *Wind*

- Energy*, 17(11), 1753–1766.
- Ospina, R., & Ferrari, S. L. (2010). Inflated beta distributions. *Statistical Papers*, 51(1), 111–126.
- Peplonska, B., Bukowska, A., Sobala, W., Reszka, E., Gromadzińska, J., Wasowicz, W., ... Ursin, G. (2012). Rotating night shift work and mammographic density. *Cancer, Epidemiology, Biomarkers, & Prevention*, 21(7), 1028–1037.
- Pereira, G. H. (2017). On quantile residuals in beta regression. *Communications in Statistics - Simulation and Computation*, 1–15.
- Pratto, F., Liu, J. H., Levin, S., Sidanius, J., Shih, M., Bachrach, H., & Hegarty, P. (2000). Social dominance orientation and the legitimization of inequality across cultures. *Journal of Cross-Cultural Psychology*, 31(3), 369–409.
- Pratto, F., Sidanius, J., Stallworth, L. M., & Malle, B. F. (1994). Social dominance orientation: A personality variable predicting social and political attitudes. *Journal of Personality and Social Psychology*, 67(4), 741–763.
- Pratto, F., Stallworth, L. M., & Conway-Lanz, S. (1998). Social dominance orientation and the ideological legitimization of social policy. *Journal of Applied Social Psychology*, 28(20), 1853–1875.
- Pratto, F., Stallworth, L. M., & Sidanius, J. (1997). The gender gap: Differences in political attitudes and social dominance orientation. *British Journal of Social Psychology*, 36(1), 49–68.
- R Core Team. (2017). *R: A language and environment for statistical computing*. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from <https://www.R-project.org/>
- Rigby, R. A., & Stasinopoulos, D. M. (2005). Generalized additive models for location, scale and shape, (with discussion). *Applied Statistics*, 54(3), 507–554.
- Rigby, R. A., Stasinopoulos, D. M., Heller, G., & De Bastiani, F. (2017). *Distributions for modelling location, scale, and shape: Using GAMLSS in R*. Unpublished draft.

Retrieved from gamlss.com/books-articles

- Rosopa, P. J., Schaffer, M. M., & Schroeder, A. N. (2013). Managing heteroscedasticity in general linear models. *Psychological Methods*, 18(3), 335–351.
- Smithson, M., & Merkle, E. C. (2013). *Generalized linear models for categorical and continuous limited dependent variables*. Boca Raton, FL: CRC Press.
- Smithson, M., & Verkuilen, J. (2006). A better lemon squeezer? Maximum-likelihood regression with beta-distributed dependent variables. *Psychological Methods*, 11(1), 54–71.
- Smyth, G. K. (1989). Generalized linear models with varying dispersion. *Journal of the Royal Statistical Society. Series B (Methodological)*, 51(1), 47–60.
- Smyth, G. K., & Verbyla, A. P. (1999). Adjusted likelihood methods for modelling dispersion in generalized linear models. *Environmetrics*, 10(6), 695–709.
- Stasinopoulos, D. M., Rigby, R. A., Heller, G. Z., Voudouris, V., & De Bastiani, F. (2017). *Flexible regression and smoothing: Using GAMLSS in R*. Boca Raton, FL: CRC Press.
- Venables, W. N., & Ripley, B. D. (2002). *Modern applied statistics with S*. New York, NY: Springer.
- White, M. H., II, & Crandall, C. S. (2017). Freedom of racist speech: Ego and expressive threats. *Journal of Personality and Social Psychology*, 113(3), 413–429.
- Wilson, M. S., & Sibley, C. G. (2013). Social dominance orientation and right-wing authoritarianism: Additive and interactive effects on political conservatism. *Political Psychology*, 34(2), 277–284.
- Wooldridge, J. M. (2010). *Econometric analysis of cross section and panel data*. Cambridge, MA: MIT Press.

Table 1

Guide for Choosing Correct Distribution and Corresponding gamlss Family

Any $y_i = 0$?	Any $y_i = 1$?	Distribution	gamlss Family	Submodels
No	No	Beta	BE()	μ, σ
Yes	No	Zero-Inflated Beta	BEINF0()	μ, σ, ν
No	Yes	One-Inflated Beta	BEINF1()	μ, σ, ν
Yes	Yes	Zero-and-One Inflated Beta	BEINF()	μ, σ, ν, τ

Table 2

Results From the Beta Regression Model

Submodel	Variable	b	SE	t	p
μ	Intercept	-2.03	0.20	-10.15	< .001
	Conservatism	0.28	0.05	5.46	< .001
σ	Intercept	-0.97	0.22	-4.42	< .001
	Conservatism	0.14	0.06	2.48	.014
ν	Intercept	0.82	0.38	2.13	.034
	Conservatism	-0.54	0.12	-4.46	< .001
τ	Intercept	-13.33	6.91	-1.93	.055
	Conservatism	1.70	1.15	1.47	.143




Figure 1. Beta probability density functions with various combinations of shape parameters.




Figure 2. The stronger the normative influence on prejudice, the smaller the variance of reported prejudice.




Figure 3. The more socially acceptable the prejudice, the less people opt for the politically correct response.




Figure 4. The distribution of SDO scores contains a floor effect.




Figure 5. The residuals plot shows heteroskedasticity and conditional non-normality.




Figure 6. SDO regressed on conservatism, using OLS regression. Dotted lines are one standard deviation (i.e., estimated standard deviation of the residuals) above and below the solid regression line.




Figure 7. SDO regressed on conservatism, using beta regression. Dotted lines are one standard deviation (i.e., estimated standard deviation of the residuals) above and below the solid regression line.